

06 Implement SGD

February 28, 2020

1 Implementing SGD

```
In [1]: import warnings
        warnings.filterwarnings("ignore")

        from sklearn.datasets import load_boston
        from random import seed
        from random import randrange
        from csv import reader
        from math import sqrt
        from sklearn import preprocessing
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        from prettytable import PrettyTable
        from sklearn.linear_model import SGDRegressor
        from sklearn import preprocessing
        from sklearn.metrics import mean_squared_error

In [22]: from sklearn.model_selection import train_test_split
         import sklearn

         from sklearn.datasets import load_boston # to load datasets from sklearn
         import matplotlib.pyplot as plt

         import numpy as np
         import seaborn as sns

         from collections import Counter
         from sklearn.metrics import accuracy_score

         from sklearn.preprocessing import StandardScaler
         import pandas as pd
         import math

In [5]: boston = load_boston()
        # Shape of Boston datasets
        print(boston.data.shape)
```

(506, 13)

```
In [10]: print(boston.DESCR)
```

Boston House Prices dataset

=====

Notes

Data Set Characteristics:

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive

:Median Value (attribute 14) is usually the target

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<http://archive.ics.uci.edu/ml/datasets/Housing>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on

pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

****References****

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity'
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the AAAI Conference on Artificial Intelligence
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>)

```
In [11]: boston.feature_names
```

```
Out[11]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [12]: columnNames = boston.feature_names  
        columnNames
```

```
Out[12]: array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
               'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

```
In [19]: boston.data
```

```
Out[19]: array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02,  
                4.9800e+00],  
               [2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,  
                9.1400e+00],  
               [2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,  
                4.0300e+00],  
               ...,  
               [6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,  
                5.6400e+00],  
               [1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,  
                6.4800e+00],  
               [4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,  
                7.8800e+00]])
```

```
In [13]: Data = pd.DataFrame(boston.data, columns = columnNames)  
        Data.head(2)
```

```
Out[13]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	

	PTRATIO	B	LSTAT
0	15.3	396.9	4.98
1	17.8	396.9	9.14

```
In [17]: Data_Labels = boston.target
        Data_Labels.shape
```

```
Out[17]: (506,)
```

```
In [15]: Data["PRICE"] = Data_Labels
        Data.head(2)
```

```
Out[15]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.9	4.98	24.0
1	17.8	396.9	9.14	21.6

```
In [20]: Data.shape
```

```
Out[20]: (506, 14)
```

2 Train Test Split

```
In [23]: X_train, X_test, Y_train, Y_test = train_test_split(Data, Data["PRICE"], test_size = 0.2
        X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
        print(X_train.head(10))
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
243	0.12757	30.0	4.93	0.0	0.428	6.393	7.8	7.0355	6.0	300.0	
253	0.36894	22.0	5.86	0.0	0.431	8.259	8.4	8.9067	7.0	330.0	
224	0.31533	0.0	6.20	0.0	0.504	8.266	78.3	2.8944	8.0	307.0	
466	3.77498	0.0	18.10	0.0	0.655	5.952	84.7	2.8715	24.0	666.0	
460	4.81213	0.0	18.10	0.0	0.713	6.701	90.0	2.5975	24.0	666.0	
55	0.01311	90.0	1.22	0.0	0.403	7.249	21.9	8.6966	5.0	226.0	
75	0.09512	0.0	12.83	0.0	0.437	6.286	45.0	4.5026	5.0	398.0	
337	0.03041	0.0	5.19	0.0	0.515	5.895	59.6	5.6150	5.0	224.0	
249	0.19073	22.0	5.86	0.0	0.431	6.718	17.5	7.8265	7.0	330.0	
485	3.67367	0.0	18.10	0.0	0.583	6.312	51.9	3.9917	24.0	666.0	

	PTRATIO	B	LSTAT	PRICE
243	16.6	374.71	5.19	23.7
253	19.1	396.90	3.54	42.8
224	17.4	385.05	4.14	44.8
466	20.2	22.01	17.15	19.0
460	20.2	255.23	16.42	16.4
55	17.9	395.93	4.81	35.4
75	18.7	383.23	8.94	21.4
337	20.2	394.81	10.56	18.5
249	19.1	393.74	6.56	26.2
485	20.2	388.62	10.58	21.2

```

In [27]: #First Standadize the data
         from sklearn import preprocessing
         scaler=preprocessing.StandardScaler()
         std_scale = scaler.fit(X_train[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'R
         train_standardized= std_scale.transform(X_train[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'R
         test_standardized= std_scale.transform(X_test[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM',

In [28]: X_train_standardized = pd.DataFrame(train_standardized, columns = columnNames)
         X_test_standardized = pd.DataFrame(test_standardized, columns = columnNames)
         Y_train=np.array(Y_train)
         Y_test=np.array(Y_test)
         print(Y_train)
         print(Y_train.shape)
         print(X_train_standardized.shape)

[23.7 42.8 44.8 19.  16.4 35.4 21.4 18.5 26.2 21.2 32.7 30.1 32.2 22.
 21.7 13.8 21.6 31.6 27.9 11.7 32.5 19.3 24.6 19.4 24.8 11.5 20.4 22.6
 14.4 19.5 19.3 28.7 31.   5.6 34.9 22.9 23.3 48.5 24.2 20.2  7.2 37.
 20.3 20.4 20.1 23.8 24.3  8.8 23.  21.9 19.3 50.  25.2 19.1 50.  19.8
 23.2 23.1 50.  50.   7.5 24.7 21.8 22.6 22.6 25.  28.4 15.  20.  20.
 18.4 50.  23.9 20.5 22.8 24.8 22.2 14.2 17.  12.6 34.7 25.1 14.9 30.8
 18.3 48.8 13.5 23.4 26.4 20.  14.5 10.5  8.4 22.9 23.1 23.3 18.9 23.
 35.2 10.5 23.3 19.5 13.2 22.2 17.6 28.2 17.8 23.9 28.1 14.3 13.4  5.
 13.8 14.  46.7 33.  24.8 16.1 22.1 33.1 16.7 18.2 10.9 20.1 16.8 21.5
 20.6 15.  20.2 23.8 13.4 19.2 12.5 25.   7.  13.8 19.4 15.4 14.6 23.1
 17.1 13.5 26.6 25.  14.9  9.7 17.5 23.2 18.5 23.2 31.6 24.1 20.6  8.3
 22.5 10.2  6.3 18.9 15.6  8.4  8.7 46.  19.5 15.  21.  20.8 13.3 32.
 12.1 28.  26.4 50.  21.6 16.1 18.3 23.1 43.1 23.  31.5 20.5 17.1 21.2
 12.7 14.5 18.6 12.  21.7 19.4 36.2 11.8 18.2 50.  50.  31.1 50.  42.3
 10.2 28.7 23.5 15.6 21.1 13.  27.  22.9 20.6 28.7 19.6 25.  22.2 50.
 34.9 20.3 14.1 15.2 22.5 18.5 13.8 20.1 38.7 20.3 30.5 13.3 23.6 37.2
 23.1 22.6 33.1 22.4 22.  14.1 22.2 20.4 25.  18.8 24.6 17.8 17.3 19.8
 22.6 13.1 15.6 14.8 17.8 20.6 33.4 26.7 21.4 23.9 18.  22.  50.  11.8
 13.6 28.5 37.3  7.2 24.5 19.2 17.4 20.8 14.6 31.5 43.8 31.2 27.1 18.9
 22.3 24.8 37.9 27.9 10.4 23.3 18.5 15.2 50.  16.2 25.  29.6 20.1 19.4
 29.9 17.5 24.4 22.9 19.7 19.4 13.1 29.8 22.  20.6 24.3 25.3 37.6 33.2
 20.  20.6 20.  24.7 17.4 19.   8.1 21.7 10.9 50.  20.4 19.1 21.7 21.8
 43.5 27.1 10.2 44.  23.6 24.4 19.7 39.8 16.1 32.9 17.1 24.1 35.1 21.2
 25.  27.5 20.3 16.7 19.6 23.9 18.7 17.2 32.  19.6 18.8 29.6 33.2 22.
 21.  19.6 17.8 19.3 13.9 36.1 19.8 36.2 34.6 23.7 50.  24.7 12.3 19.9
 33.3 20.5 13.9 20.1 20.8 23.8 21.2 21.4 25.  19.1 16.6 36.4 27.5 48.3
 16.  45.4 13.1 18.2  7.4 16.2 19.5 23.1 21.  32.4 11.  21.4 23.   5.
 20.7 29.8 29.1 21.7 14.3  9.6 19.6 13.8 16.5 22.5 23.7 22.   8.8 22.7
 33.8 14.5  8.5 21.9 11.9 12.8 36.5 22.7 15.2 19.1 21.9 22.8]
(404,)
(404, 13)

In [29]: X_train_standardized['PRICE']=Y_train

```

```
X_test_standardized['PRICE']=Y_test
```

```
In [30]: print(X_train_standardized.shape)
         print(X_test_standardized.shape)
```

```
(404, 14)
```

```
(102, 14)
```

3 Stochastic Gradient Descent

```
In [32]: #have used the below code from github
```

```
In [31]: # for references
```

```
#https://github.com/gauravtheP/Implementing-Stochastic-GradientDescent/blob/master/Li
#First step initilize the weights and b
#formulae of slope s=mx+b.
# mx is the weights*x1...weights_d*xd
# b is the intercept term
m = X_train.shape[0]
weight = np.random.randn(13)*np.sqrt(2/m) # defining initial random weight from norma
b = np.random.randn(1)*np.sqrt(2/m) # generating initial random y-intercept from norm
# initilize learing rate
learningRate = 0.2
print(m,weight,b,learningRate)
for i in range(2000): # running 2000 iterations
    Data_batch_10 = X_train_standardized.sample(n = 10) # taking 10 stochastic sample
    X_temp = Data_batch_10.drop("PRICE", axis = 1, inplace = False) # DROP the price
    #X = pd.DataFrame(X_temp, columns = columnNames)
    X=X_temp
    Y = Data_batch_10["PRICE"]
    PartialGradient = np.empty(13)# in this we store the partial derivate with respec
    sum2 = 0
    # Update the weights-----
    # formula (w0=w1-lr*derivate)in every iteration
    # step 1.
    #First calculate the derivative
    for j in range(13): # as there are 13 dimensions in our dataset and dimensions of
        sum1 = 0
        for k in range(10):
            sum1 += -2 * X.iloc[k][j] * (Y.iloc[k] - np.dot(weight, X.iloc[k]) - b) #
        PartialGradient[j] = sum1
    # step 2.
    #multiply with learning rate
    PartialGradient *= learningRate
    #step 3.
    #Update the weights
    for l in range(13):
```

```

        weight[1] -= PartialGradient[1] # updating weights
# Update the Intercepts or (b's)-----
    for m in range(10):
        sum2 += -2 * (Y.iloc[m]- np.dot(weight, X.iloc[m]) - b) # this is the derivat
        b = b - learningRate * sum2 #updating y-intercept 'b'
# in every iteration u have to reduce the learing rate bro
        learningRate = 0.01 / pow(i+1, 0.25) #learning rate at every iteration
# just add the regularization term to it
        weight = weight + 0.0001*np.dot(weight, weight) #adding l2 regularization
        b = b + 0.0001*np.dot(weight, weight) #adding L2 regularization
    print("Weight = "+str(weight))
    print("b = "+str(b))

404 [ 0.02381462  0.05560947  0.09227067  0.0094231   0.08282602 -0.09006225
      -0.04582119 -0.02474342  0.03889422  0.07088581  0.06431994  0.01697569
      -0.23645988] [-0.05827495] 0.2
Weight = [-0.77962561  1.46734951  0.57203129  1.10257263 -0.82167183  3.46647483
          0.40517137 -2.24246557  2.94903055 -2.19957943 -1.73340329  1.55374565
          -3.57664596]
b = [23.1528653]

```

```

In [41]: # time for testdata.. with our updated weights and coeffcients
import math
test_temp = X_test_standardized.drop("PRICE", axis = 1, inplace = False)
test_data = test_temp
test_labels = Y_test
y_predicted = []
for i in range(102):
    test_i = 0
    test_i = np.dot(weight, test_data.iloc[i]) + b[0] #making prediction by using opt
    y_predicted.append(test_i)

```

```

In [42]: #Make the predictions
d1 = {'True Labels': Y_test, 'Predicted Labels': y_predicted}
df1 = pd.DataFrame(data = d1)

```

```

In [43]: Mean_Sq_Error = mean_squared_error(Y_test, y_predicted)
print(Mean_Sq_Error)

```

```


19.430696876801047

```

```

In [44]: import seaborn as sns
lm1 = sns.lmplot(x="True Labels", y="Predicted Labels", data = df1, size = 10)
fig1 = lm1.fig
fig1.suptitle("With Manual Implementation(Without Sklearn)", fontsize=18)
sns.set(font_scale = 1.5)

```



```
In [45]: #Sklearn implementation
X_temp = X_train_standardized.drop("PRICE", axis = 1, inplace = False)
X=X_temp
Y = Y_train
X_test_temp = X_test_standardized.drop("PRICE", axis = 1, inplace = False)
X_te=X_test_temp
Y_te = Y_test
clf = SGDRegressor(shuffle = False, learning_rate= 'invscaling', max_iter = 2000)
clf.fit(X, Y)# fir train data
Y_pred = clf.predict(X_te)# predict test error
print("Weight = "+str(clf.coef_))
print("Y Intercept = "+str(clf.intercept_))
```



```

Weight = [-1.17297343  1.33565484  0.03873578  0.6983058  -1.52908145  2.43744545
 -0.02551971 -3.29220995  2.82299091 -2.19753749 -2.00640482  0.98792572
 -4.09702758]
Y Intercept = [22.80739133]

```

```

In [46]: d2 = {'True Labels': Y_te, 'Predicted Labels': Y_pred}
df2 = pd.DataFrame(data = d2)
df2

```

```

Out[46]:

```

	True Labels	Predicted Labels
0	23.2	26.961971
1	11.7	13.973693
2	21.5	24.903008
3	18.4	16.019471
4	21.2	20.579715
5	29.4	31.096605
6	34.9	34.723440
7	15.7	14.071281
8	30.7	31.618960
9	16.6	17.117845
10	19.4	16.610444
11	17.5	16.819046
12	23.7	28.307462
13	13.3	13.777073
14	50.0	41.876746
15	14.9	14.582086
16	35.4	34.247308
17	8.5	6.438918
18	22.3	27.041883
19	8.3	13.126012
20	21.4	20.714101
21	24.0	25.351100
22	13.1	14.044207
23	21.1	20.824257
24	13.6	13.499123
25	12.7	10.908501
26	27.5	18.833419
27	21.7	20.716709
28	26.5	25.710582
29	20.7	25.814126
..
72	23.4	23.916236
73	19.9	18.412606
74	22.8	24.373119
75	33.4	28.444640
76	16.5	10.390914
77	24.5	20.716349

78	17.8	22.434002
79	24.4	23.468390
80	28.6	27.910439
81	19.3	16.588466
82	13.4	15.977304
83	14.4	7.674054
84	19.9	19.830509
85	28.4	31.093707
86	19.9	19.675946
87	23.9	28.043596
88	29.0	31.672917
89	29.1	31.602516
90	11.9	22.966558
91	17.2	15.639889
92	41.7	37.748548
93	14.1	15.471684
94	11.3	13.725981
95	15.6	13.141447
96	30.3	32.204866
97	13.4	15.625202
98	17.9	0.238130
99	23.1	20.996016
100	50.0	35.756602
101	24.4	23.983928

[102 rows x 2 columns]

```
In [47]: Mean_Sq_Error = mean_squared_error(Y_te, Y_pred)
Mean_Sq_Error
```

```
Out[47]: 18.474840956762705
```

```
In [49]: lm2 = sns.lmplot(x="True Labels", y="Predicted Labels", data = df2, size = 10)
fig2 = lm2.fig
# Add a title to the Figure
fig2.suptitle("With Sklearn's Implementation", fontsize=18)
sns.set(font_scale = 1.5)
```

