

Model3_LSTM

May 3, 2020

1 DonorsChoose

DonorsChoose.org receives hundreds of thousands of project proposals each year for classroom projects in need of funding. Right now, a large number of volunteers is needed to manually screen each submission before it's approved to be posted on the DonorsChoose.org website.

Next year, DonorsChoose.org expects to receive close to 500,000 project proposals. As a result

How to scale current manual processes and resources to screen 500,000 projects so that they can
How to increase the consistency of project vetting across different volunteers to improve t
How to focus volunteer time on the applications that need the most assistance

The goal of the competition is to predict whether or not a DonorsChoose.org project proposal submitted by a teacher will be approved, using the text of project descriptions as well as additional metadata about the project, teacher, and school. DonorsChoose.org can then use this information to identify projects most likely to need further review before approval.

1.1 About the DonorsChoose Data Set

The `train.csv` data set provided by DonorsChoose contains the following features:

Feature	Description
<code>project_id</code>	A unique identifier for the proposed project. Example: p036502

`project_title` | Title of the project. **Examples:**

Art Will Make You Happy!

First Grade Fun

`project_grade_category` | Grade level of students for which the project is targeted. One of the following enumerated values:

Grades PreK-2

Grades 3-5

Grades 6-8

Grades 9-12

`project_subject_categories` | One or more (comma-separated) subject categories for the project from the following enumerated list of values:

Applied Learning

Care & Hunger
 Health & Sports
 History & Civics
 Literacy & Language
 Math & Science
 Music & The Arts
 Special Needs
 Warmth

Examples:

Music & The Arts
 Literacy & Language, Math & Science

school_state | State where school is located ([Two-letter U.S. postal code](#)). **Example:** WY
project_subject_subcategories | One or more (comma-separated) subject subcategories for the project. **Examples:**

Literacy
 Literature & Writing, Social Sciences

project_resource_summary | An explanation of the resources needed for the project. **Example:**

My students need hands on literacy materials to manage sensory needs!

project_essay_1 | First application essay

project_essay_2 | *Second application essay* **project_essay_3** | Third application essay

project_essay_4 | *Fourth application essay* **project_submitted_datetime** | Datetime when project application was submitted. **Example:** 2016-04-28 12:43:56.245

teacher_id | A unique identifier for the teacher of the proposed project. **Example:** bdf8baa8fedef6bfeec7ae4ff1c15c56

teacher_prefix | Teacher's title. One of the following enumerated values:

nan
 Dr.
 Mr.
 Mrs.
 Ms.
 Teacher.

teacher_number_of_previously_posted_projects | Number of project applications previously submitted by the same teacher. **Example:** 2

* See the section Notes on the Essay Data for more details about these features.

Additionally, the `resources.csv` data set provides more data about the resources required for each project. Each line in this file represents a resource required by a project:

Feature	Description
id	A <code>project_id</code> value from the <code>train.csv</code> file. Example: p036502
description	Description of the resource. Example: Tenor Saxophone Reeds, Box of 25

Feature	Description
quantity	Quantity of the resource required.
price	Price of the resource required. Example: 9.95

Note: Many projects require multiple resources. The `id` value corresponds to a `project_id` in `train.csv`, so you use it as a key to retrieve all resources needed for a project:

The data set contains the following label (the value you will attempt to predict):

Label	Description
project_is_approved	Advisory flag indicating whether DonorsChoose approved the project. A value of 0 indicates the project was not approved, and a value of 1 indicates the project was approved.

1.1.1 Notes on the Essay Data

Prior to May 17, 2016, the prompts for the essays were as follows:

project_essay_1: "Introduce us to your classroom"

project_essay_2: "Tell us more about your students"

project_essay_3: "Describe how your students will use the materials you're requesting"

project_essay_3: "Close by sharing why your project will make a difference"

Starting on May 17, 2016, the number of essays was reduced from 4 to 2, and the prompts for the first 2 essays were changed to the following:

project_essay_1: "Describe your students: What makes your students special? Specific details about their background, your neighborhood, and your school are all helpful."

project_essay_2: "About your project: How will these materials make a difference in your students' learning and improve their school lives?"

For all projects with `project_submitted_datetime` of 2016-05-17 and later, the values of `project_essay_3` and `project_essay_4` will be NaN.

```
In [0]: !pip install chart_studio
```

```
Collecting chart_studio
```

```
  Downloading https://files.pythonhosted.org/packages/ca/ce/330794a6b6ca4b9182c38fc69dd2a9cbff/
    || 71kB 2.3MB/s
```

```
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from chart_studio)
```

Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from chart_studio)

Installing collected packages: chart-studio

Successfully installed chart-studio-1.1.0

```
In [0]: import chart_studio.plotly as py
import plotly.graph_objs as go
```

```
In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os
```

```

#from plotly import plotly
import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

```

In [0]: %tensorflow_version 1.x

TensorFlow 1.x selected.

```

In [0]: # Load the Drive helper and mount
        from google.colab import drive

        # This will prompt for authorization.
        drive.mount('/content/drive')

```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-

Enter your authorization code:

uuuuuuuuuuuu

Mounted at /content/drive

1.2 1.1 Reading Data

```

In [0]: resource_data=pd.read_csv('/content/drive/My Drive/LSTM Assignment/resources.csv')
        project_data=pd.read_csv('/content/drive/My Drive/LSTM Assignment/train_data.csv')
        #project_data=project_data.sample(n=1000)

```

```

In [0]: price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
        project_data = pd.merge(project_data, price_data, on='id', how='left')

```

1.3 1.2 preprocessing of project_subject_categories

```

In [0]: categories = list(project_data['project_subject_categories'].values)
        # remove special characters from list of strings python: https://stackoverflow.com/a/4758804

        # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
        # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
        # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
        cat_list = []
        for i in categories:
            temp = ""
            # consider we have text like this "Math & Science, Warmth, Care & Hunger"
            for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
                if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
                    j=j.replace('The','') # if we have the words "The" we are going to replace them with ''
                j = j.replace(' ','') # we are replacing all the ' '(space) with ''(empty) ex: "Math & Science" becomes "Math&Science"
            cat_list.append(temp+j)

```

```

        temp+=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing sp
        temp = temp.replace('&','_') # we are replacing the & value into
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

1.4 1.3 preprocessing of project_subject_subcategories

```

In [0]: sub_categories = list(project_data['project_subject_subcategories'].values)
        # remove special characters from list of strings python: https://stackoverflow.com/a/4

        # https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
        # https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-str
        # https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-pyt

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth
        if 'The' in j.split(): # this will split each of the category based on space "
            j=j.replace('The','') # if we have the words "The" we are going to replace
        j = j.replace(' ','') # we are placing all the ' '(space) with ''(empty) ex: "
        temp +=j.strip()+" " #" abc ".strip() will return "abc", remove the trailing sp
        temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/408403
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

In [0]: # https://stackoverflow.com/a/47091490/4084039

```

```

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

```

```

In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're",
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'I',
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that's",
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no',
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won',
            "won't", 'wouldn', "wouldn't"]

```

```

In [0]: # Combining all the above students
def Text_cleaner(data):

    preprocessed_essays = []
    # tqdm is for printing the status bar
    for sentence in tqdm(data.values):
        sent = decontracted(sentence)
        sent = sent.replace('\r', ' ')
        sent = sent.replace('\n', ' ')
        sent = sent.replace('\n', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280

```

```

        sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
        preprocessed_essays.append(sent.lower().strip())
    return preprocessed_essays

```

In [0]: # merge two column text dataframe:

```

project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)

```

In [0]: project_data['essay']=Text_cleaner(project_data['essay'])

100%| 109248/109248 [00:50<00:00, 2144.77it/s]

In [0]: #<https://stackoverflow.com/questions/29763620/how-to-select-all-columns-except-one-col>

```

X=project_data.loc[:, project_data.columns != 'project_is_approved']
y=project_data['project_is_approved']
X.shape

```

Out[0]: (109248, 19)

In [0]: X.head()

```

Out[0]:
          essay  ... quantity
0  students english learners working english seco...  ...         23
1  students arrive school eager learn polite gene...  ...          1
2  true champions not always ones win guts mia ha...  ...         22
3  work unique school filled esl english second l...  ...          4
4  second grade classroom next year made around 2...  ...          4

```

[5 rows x 9 columns]

1.5 Splitting data into Test,Train,CV

In [0]: from sklearn.model_selection import train_test_split

```

X_train,X_test,y_train, y_test=train_test_split(X, y, test_size=0.2)
X_train,X_cv,y_train,y_cv=train_test_split(X_train, y_train, test_size=0.2)
print(X_train.shape)
print(X_test.shape)

print(y_train.shape)
print(y_test.shape)

```

(69918, 9)

(21850, 9)

(69918,)

(21850,)

1.5.1 Defining a custom metric AUC

2 Model-3

ref: <https://i.imgur.com/fkQ8nGo.png>

- **input_seq_total_text_data:**
- **Other_than_text_data:**
 - . Convert all your Categorical values to onehot coded and then concatenate all these onehot vectors . Neumerical values and use CNN1D as shown in above figure. . You are free to choose all CNN parameters like kernel sizes, stride.

2.1 Load the GloVe from text

```
In [0]: from numpy import asarray
def get_glove():
    embeddings_index = dict()
    f = open('/content/glove.6B.300d.txt')
    for line in tqdm(f):
        values = line.split()
        word = values[0]
        coefs = asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs
    f.close()
    return embeddings_index
```

```
In [0]: embeddings_index=get_glove()
```

400000it [00:27, 14493.17it/s]

```
In [0]: from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from numpy import zeros
```

```
#def word_embedding(docs, embeddings_index):
glove_words = set(embeddings_index.keys())

sentence_vectors = []
# prepare tokenizer
t = Tokenizer()
t.fit_on_texts(X_train['essay'].values)
vocab_size = len(t.word_index) + 1
# integer encode the documents
X_train_encoded_docs = t.texts_to_sequences(X_train['essay'].values)
X_test_encoded_docs = t.texts_to_sequences(X_test['essay'].values)
X_cv_encoded_docs = t.texts_to_sequences(X_cv['essay'].values)
```

```

# pad documents to a max length of 300 words
max_length = 300
X_train_padded_docs = pad_sequences(X_train_encoded_docs, maxlen=max_length, padding='post')
X_test_padded_docs = pad_sequences(X_test_encoded_docs, maxlen=max_length, padding='post')
X_cv_padded_docs = pad_sequences(X_cv_encoded_docs, maxlen=max_length, padding='post')

embedding_matrix = zeros((vocab_size, max_length))
for word, i in t.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector

print(X_train_padded_docs.shape)
print(X_test_padded_docs.shape)
print(X_cv_padded_docs.shape)

```

Using TensorFlow backend.

```

(69918, 300)
(21850, 300)
(17480, 300)

```

```

In [0]: essay_input = Input(shape=(len(X_train_padded_docs[0]),), name='essay_input')
        essay_input_1 = Embedding(input_dim=vocab_size, output_dim=300, input_length=len(X_train_padded_docs[0]))
        essay_input_1 = LSTM(32, return_sequences=True)(essay_input_1)
        essay_input_1 = Flatten()(essay_input_1)

```

2.1.1 Vectorize the Features

2.2 Categorical data: clean_categories

2.2.1 Building train, test and validation data

```

In [0]: # we use count vectorizer to convert the values into one
        from sklearn.feature_extraction.text import CountVectorizer
        vectorizer = CountVectorizer(lowercase=False, binary=True)
        vectorizer.fit(X['clean_categories'].values)

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_cat=vectorizer.transform(X_train['clean_categories'].values)
        X_test_cat=vectorizer.transform(X_test['clean_categories'].values)
        X_cv_cat=vectorizer.transform(X_cv['clean_categories'].values)

        print(vectorizer.get_feature_names())
        print(X_train_cat.shape)

```

```
print(X_test_cat.shape)
print(X_cv_cat.shape)
```

```
['AppliedLearning', 'Care_Hunger', 'Health_Sports', 'History_Civics', 'Literacy_Language', 'Ma
(69918, 9)
(21850, 9)
(17480, 9)
```

```
In [0]: categories_input = Input(shape=(X_train_cat.shape[1],), name='categories_input')
        categories_input_1= Embedding(input_dim=X_train_cat.shape[1],output_dim=64, input_length=
        categories_input_1 = Flatten()(categories_input_1)
```

2.3 Categorical data: clean_subcategories

2.3.1 Building train, test and validation data

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
        vectorizer = CountVectorizer(lowercase=False, binary=True)
        vectorizer.fit(X['clean_subcategories'].values)

        # we use the fitted CountVectorizer to convert the text to vector
        X_train_sub_cat=vectorizer.transform(X_train['clean_subcategories'].values)
        X_test_sub_cat=vectorizer.transform(X_test['clean_subcategories'].values)
        X_cv_sub_cat=vectorizer.transform(X_cv['clean_subcategories'].values)

        print(vectorizer.get_feature_names())
        print(X_train_sub_cat.shape)
        print(X_test_sub_cat.shape)
        print(X_cv_sub_cat.shape)
```

```
['AppliedSciences', 'Care_Hunger', 'CharacterEducation', 'Civics_Government', 'College_CareerP
(69918, 30)
(21850, 30)
(17480, 30)
```

```
In [0]: sub_categories_input = Input(shape=(X_train_sub_cat.shape[1],), name='sub_categories_in
        sub_categories_input_1= Embedding(input_dim=X_train_sub_cat.shape[1],output_dim=64, in
        sub_categories_input_1 = Flatten()(sub_categories_input_1)
```

2.4 Categorical data: project_grade_category

2.4.1 Building train, test and validation data

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
        vectorizer = CountVectorizer(lowercase=False, binary=True)
        vectorizer.fit(X['project_grade_category'].values)

        # we use the fitted CountVectorizer to convert the text to vector
```

```

X_train_pro_grd=vectorizer.transform(X_train['project_grade_category'].values)
X_test_pro_grd=vectorizer.transform(X_test['project_grade_category'].values)
X_cv_pro_grd=vectorizer.transform(X_cv['project_grade_category'].values)

print(vectorizer.get_feature_names())
print(X_train_pro_grd.shape)
print(X_test_pro_grd.shape)
print(X_cv_pro_grd.shape)

```

```

['Grades35', 'Grades68', 'Grades912', 'GradesPreK2']
(69918, 4)
(21850, 4)
(17480, 4)

```

```

In [0]: proj_grade_input = Input(shape=(X_train_pro_grd.shape[1],), name='proj_grade_input')
        proj_grade_input_1= Embedding(input_dim=X_train_pro_grd.shape[1],output_dim=64, input_
        proj_grade_input_1 = Flatten()(proj_grade_input_1)

```

2.5 Categorical data: school_state

2.5.1 Building train, test and validation data

```

In [0]: from sklearn.feature_extraction.text import CountVectorizer
        vectorizer = CountVectorizer(lowercase=False, binary=True)
        vectorizer.fit(X['school_state'].values)

```

```

# we use the fitted CountVectorizer to convert the text to vector
X_train_skl_st=vectorizer.transform(X_train['school_state'].values)
X_test_skl_st=vectorizer.transform(X_test['school_state'].values)
X_cv_skl_st=vectorizer.transform(X_cv['school_state'].values)

print(vectorizer.get_feature_names())
print(X_train_skl_st.shape)
print(X_test_skl_st.shape)
print(X_cv_skl_st.shape)

```

```

['AK', 'AL', 'AR', 'AZ', 'CA', 'CO', 'CT', 'DC', 'DE', 'FL', 'GA', 'HI', 'IA', 'ID', 'IL', 'IN
(69918, 51)
(21850, 51)
(17480, 51)

```

```

In [0]: school_state_input = Input(shape=(X_train_skl_st.shape[1],), name='school_state_input')
        school_state_input_1= Embedding(input_dim=X_train_skl_st.shape[1],output_dim=64, input_
        school_state_input_1 = Flatten()(school_state_input_1)

```

2.6 Categorical data: teacher_prefix

2.6.1 Building train, test and validation data

```
In [0]: from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(lowercase=False, binary=True)
vectorizer.fit(X['teacher_prefix'].values)

# we use the fitted CountVectorizer to convert the text to vector
X_train_tch=vectorizer.transform(X_train['teacher_prefix'].values)
X_test_tch=vectorizer.transform(X_test['teacher_prefix'].values)
X_cv_tch=vectorizer.transform(X_cv['teacher_prefix'].values)

print(vectorizer.get_feature_names())
print(X_train_tch.shape)
print(X_test_tch.shape)
print(X_cv_tch.shape)

['Dr', 'Mr', 'Mrs', 'Ms', 'Teacher']
(69918, 5)
(21850, 5)
(17480, 5)

In [0]: tch_input = Input(shape=(X_train_tch.shape[1],), name='tch_input')
tch_input_1= Embedding(input_dim=X_train_tch.shape[1],output_dim=64, input_length=X_train_tch.shape[1])
tch_input_1 = Flatten()(tch_input_1)
```

2.7 price

2.7.1 Building train, test and validation data

```
In [0]: from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
scaler.fit(X_train['price'].values.reshape(-1,1))
X_train_price=scaler.transform(X_train['price'].values.reshape(-1,1))
X_test_price=scaler.transform(X_test['price'].values.reshape(-1,1))
X_cv_price=scaler.transform(X_cv['price'].values.reshape(-1,1))

print(X_train_price.shape)
print(X_test_price.shape)
print(X_cv_price.shape)

(69918, 1)
(21850, 1)
(17480, 1)
```

2.8 quantity

2.8.1 Building train and validation data

```
In [0]: scaler = MinMaxScaler()
        scaler.fit(X_train['quantity'].values.reshape(-1,1))
        X_train_quantity=scaler.transform(X_train['quantity'].values.reshape(-1,1))
        X_test_quantity=scaler.transform(X_test['quantity'].values.reshape(-1,1))
        X_cv_quantity=scaler.transform(X_cv['quantity'].values.reshape(-1,1))

        print(X_train_quantity.shape)
        print(X_test_quantity.shape)
        print(X_cv_quantity.shape)

(69918, 1)
(21850, 1)
(17480, 1)
```

2.9 teacher_number_of_previously_posted_projects

2.9.1 Building train, test and validation data

```
In [0]: scaler = MinMaxScaler()
        scaler.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
        X_train_pre_proj=scaler.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
        X_test_pre_proj=scaler.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
        X_cv_pre_proj=scaler.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))

        print(X_train_pre_proj.shape)
        print(X_test_pre_proj.shape)
        print(X_cv_pre_proj.shape)

(69918, 1)
(21850, 1)
(17480, 1)
```



```
In [0]: # merge two sparse matrices: https://stackoverflow.com/a/19710648/4084039
        from scipy.sparse import hstack

        X_train_other = hstack((X_train_cat,X_train_sub_cat,X_train_pro_grd,X_train_skl_st,X_train_tch,X_train_price))
        X_test_other = hstack((X_test_cat,X_test_sub_cat,X_test_pro_grd,X_test_skl_st,X_test_tch,X_test_price))
        X_cv_other = hstack((X_cv_cat,X_cv_sub_cat,X_cv_pro_grd,X_cv_skl_st,X_cv_tch,X_cv_price))

        print(X_train_other.shape)
        print(X_test_other.shape)
        print(X_cv_other.shape)
```

```
(69918, 102)
(21850, 102)
(17480, 102)
```

```
In [0]: numeral_input=Input(shape=(X_train_numerals.shape[1],),name='numeral_input')
        numeral_input_1 = Dense(input_dim=3,output_dim=128)(numeral_input)
```

2.10 Concatenate

```
In [0]: x = concatenate([essay_input_1,other_input_1])
```

2.10.1 Defining a custom metric AUC

```
In [0]: import tensorflow as tf
        from sklearn.metrics import roc_auc_score

        def auROC(y_true, y_pred):
            try:
                return tf.py_func(roc_auc_score, (y_true, y_pred), tf.double)

            except ValueError:
                pass
```

2.11 Applying the Best found model

```
In [0]: # !pip install -q tf-nightly-2.0-preview
        # if you want to use the tf2.0 please uncomment the above line
        # Load the TensorBoard notebook extension
        # there are other ways of doing this: https://www.dlology.com/blog/quick-guide-to-run-
        %load_ext tensorboard
```

```
In [0]: # Clear any logs from previous runs
        !rm -rf ./logs/
```

```
In [0]: import tensorflow as tf
        from time import time
        import datetime
        from tensorflow.python.keras.callbacks import TensorBoard, EarlyStopping
        from keras.callbacks import TensorBoard
        from keras.callbacks import ModelCheckpoint
        #tensorboard = TensorBoard(log_dir="logs/{}".format(time))
        log_dir="logs\\fit\\" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
        tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=
        filepath="model2.h5"
        checkpoint = ModelCheckpoint(filepath, monitor='val_auroc', verbose=1, save_best_only=
```

```
In [0]: from keras.layers import Input, Embedding, LSTM, Dense,Flatten,concatenate,Dropout,LSTM
        from keras.models import Model
```

```

from keras.callbacks import EarlyStopping

# And finally we add the main logistic regression layer
output =Dense(128, activation='relu')(x)
output =Dropout(0.5)(output)
output =Dense(64, activation='relu')(output)
output =Dropout(0.5)(output)
output =Dense(32, activation='relu')(output)
output =Dense(1, activation='sigmoid',name='output')(output)

#defines model with multiple input one output
model = Model(inputs=[essay_input, other_input], outputs=output)

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy',auroc])

es = EarlyStopping(monitor='val_loss', patience=1,verbose=0, mode='min')

model.fit([X_train_padded_docs, X_train_other.todense()],y_train,
          epochs=4,
          batch_size=1000,
          validation_data=([X_cv_padded_docs,X_cv_other.todense()],y_cv),
          callbacks=[es,tensorboard_callback,checkpoint])

```

Train on 69918 samples, validate on 17480 samples

Epoch 1/4

69918/69918 [=====] - 23s 327us/step - loss: 0.4250 - accuracy: 0.8449

Epoch 00001: val_auroc did not improve from 0.71596

Epoch 2/4

69918/69918 [=====] - 22s 310us/step - loss: 0.3956 - accuracy: 0.8489

Epoch 00002: val_auroc did not improve from 0.71596

Epoch 3/4

69918/69918 [=====] - 22s 310us/step - loss: 0.3826 - accuracy: 0.8499

Epoch 00003: val_auroc did not improve from 0.71596

Epoch 4/4

69918/69918 [=====] - 22s 314us/step - loss: 0.3667 - accuracy: 0.8549

Epoch 00004: val_auroc did not improve from 0.71596

Out[0]: <keras.callbacks.callbacks.History at 0x7f6b14893780>

In [0]: *##tensorboard --logdir .*

%tensorboard --logdir logs\\fit\\20200502-200320

Output hidden; open in <https://colab.research.google.com> to view.


```
In [1]: from IPython.display import Image
        from IPython.core.display import HTML
        Image(url= "https://imgur.com/a/xxI4R0n.jpg")
```

```
Out[1]: <IPython.core.display.Image object>
```