

# 4\_DonorsChoose\_NB

March 15, 2020

## 1 Assignment 6: Apply NB

```
<li><strong>Apply Multinomial NB on these feature sets</strong>
  <ul>
    <li><font color='red'>Set 1</font>: categorical, numerical features + preprocessed_eas
    <li><font color='red'>Set 2</font>: categorical, numerical features + preprocessed_eas
  </ul>
</li>
<li><strong>The hyper paramter tuning(find best alpha:smoothing parameter)</strong>
  <ul>
<li>Find the best hyper parameter which will give the maximum <a href='https://www.appliedaicom
<li>find the best hyper paramter using k-fold cross validation(use GridsearchCV or Randomsearch
<li></li>
  </ul>
</li>
<li>
<strong>Representation of results</strong>
  <ul>
<li>You need to plot the performance of model both on train data and cross validation data for
<img src='https://i.imgur.com/hUv6aEy.jpg' width=300px></li>
<li>Once after you found the best hyper parameter, you need to train your model with it, and f
<img src='https://i.imgur.com/wMQDTFe.jpg' width=300px></li>
<li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
<img src='https://i.imgur.com/IdN5Ctv.png' width=300px></li>
  </ul>
</li>
<li>
```

fine the top 20 features from either from feature Set 1 or feature Set 2 using absolute values of feature\_log\_prob\_ parameter of MultinomialNB ([https://scikit-learn.org/stable/modules/generated/sklearn.naive\\_bayes.MultinomialNB.html](https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html)) and print their corresponding feature names

```
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
  <img src='http://i.imgur.com/YVpIGGE.jpg' width=400px>
</li>
```

## 2. Naive Bayes

## 1.1 Loading Data

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
from collections import Counter

In [2]: from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-)

Enter your authorization code:

uuuuuuuuuu

Mounted at /content/drive

```
In [0]: import pandas
#data = pandas.read_csv(r'C:\Users\ASUS\Downloads\Applied AI\Assignments - Applied AI\
project_data = pd.read_csv('/content/drive/My Drive/Assignments_DonorsChoose_2018/train
resource_data = pd.read_csv('/content/drive/My Drive/Assignments_DonorsChoose_2018/res
```

```
In [23]: # how to replace elements in list python: https://stackoverflow.com/a/2582163/4084039
cols = ['Date' if x=='project_submitted_datetime' else x for x in list(project_data.c
```

```
#sort dataframe based on time pandas python: https://stackoverflow.com/a/49702492/4084039
project_data['Date'] = pd.to_datetime(project_data['project_submitted_datetime'])
project_data.drop('project_submitted_datetime', axis=1, inplace=True)
project_data.sort_values(by=['Date'], inplace=True)
```

```
# how to reorder columns pandas python: https://stackoverflow.com/a/13148611/4084039
project_data = project_data[cols]
```

```
project_data.head(2)
```

```
Out[23]:      Unnamed: 0  ... project_is_approved
55660      8393  ...      1
76127     37728  ...      1
```

```
[2 rows x 17 columns]
```

```
In [24]: print("Number of data points in train data", resource_data.shape)
print(resource_data.columns.values)
resource_data.head(2)
```

```
Number of data points in train data (1541272, 4)
['id' 'description' 'quantity' 'price']
```

```
Out[24]:      id      description  quantity  price
0  p233245  LC652 - Lakeshore Double-Space Mobile Drying Rack      1  149.00
1  p069063      Bouncy Bands for Desks (Blue support pipes)      3   14.95
```

## 1.2 1.2 preprocessing of project\_subject\_categories

```
In [0]: categories = list(project_data['project_subject_categories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/4
# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-str
```

```

# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python
cat_list = []
for i in categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace them with ''
        j = j.replace(' ', '') # we are replacing all the ' ' (space) with '' (empty) ex: "Math & Science"
        temp+=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&','_') # we are replacing the & value into _
    cat_list.append(temp.strip())

project_data['clean_categories'] = cat_list
project_data.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in project_data['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

```

### 1.3 preprocessing of project\_subject\_subcategories

```

In [0]: sub_categories = list(project_data['project_subject_subcategories'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/4000724

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-string
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-python

sub_cat_list = []
for i in sub_categories:
    temp = ""
    # consider we have text like this "Math & Science, Warmth, Care & Hunger"
    for j in i.split(','): # it will split it in three parts ["Math & Science", "Warmth", "Care & Hunger"]
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace them with ''
        j = j.replace(' ', '') # we are replacing all the ' ' (space) with '' (empty) ex: "Math & Science"
        temp +=j.strip()+" " # " abc ".strip() will return "abc", remove the trailing spaces
    temp = temp.replace('&','_')
    sub_cat_list.append(temp.strip())

project_data['clean_subcategories'] = sub_cat_list
project_data.drop(['project_subject_subcategories'], axis=1, inplace=True)

```

```

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/408403
my_counter = Counter()
for word in project_data['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

## 1.4 Preprocessing of project\_grade\_category

```

In [0]: project_grade = list(project_data['project_grade_category'].values)
# remove special characters from list of strings python: https://stackoverflow.com/a/4

# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
# https://stackoverflow.com/questions/23669024/how-to-strip-a-specific-word-from-a-str
# https://stackoverflow.com/questions/8270092/remove-all-whitespace-in-a-string-in-pyt

grade_cat_list = []
for i in project_grade:
    # consider we have text like this:
    for j in i.split(' '): # # split by spae
        j=j.replace('Grades','')# clean grades from the row
    grade_cat_list.append(j.strip())

project_data['grade_cat_list'] = grade_cat_list
project_data.drop(['project_grade_category'], axis=1, inplace=True)

my_counter = Counter()
for word in project_data['grade_cat_list'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

## 1.5 Join train & Resource dataset

```

In [28]: price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset
project_data = pd.merge(project_data, price_data, on='id', how='left')
project_data.head(2)

```

```

Out[28]:   Unnamed: 0      id  ...  price quantity
0         8393  p205479  ...   725.05         4
1         37728  p043609  ...   213.03         8

[2 rows x 19 columns]

```

## 1.6 Train Test split

```
In [42]: y = project_data['project_is_approved'].values
        X = project_data.drop(['project_is_approved'], axis=1)
        X.head(1)
```

```
Out[42]: Unnamed: 0    ...          essay
0      8393    ...  I have been fortunate enough to use the Fairy ...

[1 rows x 19 columns]
```

```
In [0]: # train test split
        from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y)
```

## 1.7 1.3 Text preprocessing

```
In [0]: # merge two column text dataframe:
        project_data["essay"] = project_data["project_essay_1"].map(str) + \
                                project_data["project_essay_2"].map(str) + \
                                project_data["project_essay_3"].map(str) + \
                                project_data["project_essay_4"].map(str)
```

```
In [32]: project_data.head(2)
```

```
Out[32]: Unnamed: 0    ...          essay
0      8393    ...  I have been fortunate enough to use the Fairy ...
1     37728    ...  Imagine being 8-9 years old. You're in your th...

[2 rows x 20 columns]
```

```
In [33]: # printing some random reviews
        print(project_data['essay'].values[0])
        print("="*50)
        print(project_data['essay'].values[150])
        print("="*50)
        print(project_data['essay'].values[1000])
        print("="*50)
        print(project_data['essay'].values[20000])
        print("="*50)
        print(project_data['essay'].values[99999])
        print("="*50)
```

```
I have been fortunate enough to use the Fairy Tale STEM kits in my classroom as well as the ST
=====
I teach high school English to students with learning and behavioral disabilities. My students
=====
\"Life moves pretty fast. If you don't stop and look around once in awhile, you could miss it.\"
=====
```

```
\nA person's a person, no matter how small.\" (Dr.Seuss) I teach the smallest students with the
=====
My classroom consists of twenty-two amazing sixth graders from different cultures and backgrounds
=====
```

```
In [0]: # https://stackoverflow.com/a/47091490/4084039
import re
```

```
def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\'s", " is", phrase)
    phrase = re.sub(r"\'d", " would", phrase)
    phrase = re.sub(r"\'ll", " will", phrase)
    phrase = re.sub(r"\'t", " not", phrase)
    phrase = re.sub(r"\'ve", " have", phrase)
    phrase = re.sub(r"\'m", " am", phrase)
    return phrase
```

```
In [35]: sent = decontracted(project_data['essay'].values[20000])
print(sent)
print("="*50)
```

```
\nA person is a person, no matter how small.\" (Dr.Seuss) I teach the smallest students with the
=====
```

```
In [36]: # \r \n \t remove from string python: http://texthandler.com/info/remove-line-breaks-
sent = sent.replace('\r', ' ')
sent = sent.replace('\n', ' ')
sent = sent.replace('\t', ' ')
print(sent)
```

A person is a person, no matter how small. (Dr.Seuss) I teach the smallest students with the

```
In [37]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
print(sent)
```

A person is a person no matter how small Dr Seuss I teach the smallest students with the bigger

```
In [0]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're",
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'theirs',
            'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'those', 'these',
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until',
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very',
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'no',
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
            "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'won',
            "won't", 'wouldn', "wouldn't"]
```

```
In [40]: # Combining all the above students
from tqdm import tqdm
preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(project_data['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    preprocessed_essays.append(sent.lower().strip())
```

100%|| 109248/109248 [01:01<00:00, 1763.12it/s]

```
In [44]: # Combining all the above students
from tqdm import tqdm
train_preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(X_train['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    train_preprocessed_essays.append(sent.lower().strip())
```



100%|| 87398/87398 [00:49<00:00, 1778.79it/s]

```
In [45]: # Combining all the above students
from tqdm import tqdm
test_preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(X_test['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\\"', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e.lower() not in stopwords)
    test_preprocessed_essays.append(sent.lower().strip())
```

100%|| 21850/21850 [00:12<00:00, 1776.15it/s]

```
In [46]: # after preprocessing
test_preprocessed_essays[20000]
```

Out[46]: 'students world potential hands journey begins opening pages book walk door every day'

### 1.7.1 1.5.1 Vectorizing Categorical data

```
In [47]: # we use count vectorizer to convert the values into one
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(vocabulary=list(sorted_cat_dict.keys()), lowercase=False)
categories_one_hot = vectorizer.fit_transform(project_data['clean_categories'].values)
#print(vectorizer.get_feature_names())
print("Shape of matrix after one hot encoding ",categories_one_hot.shape)

cat_features = vectorizer.get_feature_names()
print(cat_features)
print(len(cat_features))
```

Shape of matrix after one hot encoding (109248, 9)

['Warmth', 'Care\_Hunger', 'History\_Civics', 'Music\_Arts', 'AppliedLearning', 'SpecialNeeds', '9

```
In [48]: # we use count vectorizer to convert the values into one
vectorizer = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=False)
sub_categories_one_hot = vectorizer.fit_transform(project_data['clean_subcategories'].values)
#print(vectorizer.get_feature_names())
print("Shape of matrix after one hot encoding ",sub_categories_one_hot.shape)
```

```

sub_cat_features=vectorizer.get_feature_names()
print(sub_cat_features)
print(len(sub_cat_features))

```

Shape of matrix after one hot encodig (109248, 30)

```

['Economics', 'CommunityService', 'FinancialLiteracy', 'ParentInvolvement', 'Extracurricular',
30

```

```

In [0]: # count of all the words in corpus python: https://stackoverflow.com/a/22898595/408403
my_counter = Counter()
for word in project_data['school_state'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

```

In [50]: # we use count vectorizer to convert the values into one
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=False)
categories_one_hot = vectorizer.fit_transform(project_data['school_state'].values)
#print(vectorizer.get_feature_names())
print("Shape of matrix after one hot encodig ",categories_one_hot.shape)

state_features=vectorizer.get_feature_names()
print(state_features)
print(len(state_features))

```

Shape of matrix after one hot encodig (109248, 51)

```

['VT', 'WY', 'ND', 'MT', 'RI', 'SD', 'NE', 'DE', 'AK', 'NH', 'WV', 'ME', 'HI', 'DC', 'NM', 'KS
51

```

```

In [0]: # count of all the words in corpus python: https://stackoverflow.com/a/22898595/408403
project_data['teacher_prefix'] = project_data['teacher_prefix'].fillna(" ")
my_counter = Counter()
for word in project_data['teacher_prefix'].values.astype('str'): #https://stackoverflowfl
    my_counter.update(word.split())

sub_fix_dict = dict(my_counter)
sorted_fix_dict = dict(sorted(sub_fix_dict.items(), key=lambda kv: kv[1]))

```

```

In [76]: # we use count vectorizer to convert the values into one
from sklearn.feature_extraction.text import CountVectorizer
vectorizer = CountVectorizer(vocabulary=list(sorted_fix_dict.keys()), lowercase=False)
prefix_one_hot = vectorizer.fit_transform(project_data['teacher_prefix'].values.astype('str'))
#print(vectorizer.get_feature_names())
print("Shape of matrix after one hot encodig ",prefix_one_hot.shape)

```

```

prefix_features=vectorizer.get_feature_names()
print(prefix_features)
print(len(prefix_features))

```

Shape of matrix after one hot encodig (109248, 5)

```
['Dr.', 'Teacher', 'Mr.', 'Ms.', 'Mrs.']
```

5

In [0]: *# count of all the words in corpus python: <https://stackoverflow.com/a/22898595/408403>*

```

my_counter = Counter()
for word in project_data['grade_cat_list'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

```

In [67]: *# we use count vectorizer to convert the values into one*

```

from sklearn.feature_extraction.text import CountVectorizer
vectorizer1 = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=False)
grade_one_hot = vectorizer1.fit_transform(project_data['grade_cat_list'].values)
print(vectorizer1.get_feature_names())
print("Shape of matrix after one hot encodig ", grade_one_hot.shape)

```

```
['9-12', '6-8', '3-5', 'PreK-2']
```

Shape of matrix after one hot encodig (109248, 4)

## 1.7.2 1.5.2 Vectorizing Text data

### 1.5.2.1 Bag of words

In [55]: *# We are considering only the words which appeared in at least 10 documents(rows or p*

```

vectorizer = CountVectorizer(min_df=10)
text_bow = vectorizer.fit_transform(preprocessed_essays)
print("Shape of matrix after one hot encodig ", text_bow.shape)

essays_features=vectorizer.get_feature_names()
#print(essays_features)
print(len(essays_features))

```

Shape of matrix after one hot encodig (109248, 16512)

16512

### 1.5.2.2 TFIDF vectorizer

```

In [99]: from sklearn.feature_extraction.text import TfidfVectorizer
vectorizer3 = TfidfVectorizer(min_df=10)

```

```

text_tfidf = vectorizer3.fit_transform(preprocessed_essays)
print("Shape of matrix after one hot encoding ",text_tfidf.shape)

essays_features_tf=vectorizer3.get_feature_names()
#print(essays_features)
print(len(essays_features_tf))

```

Shape of matrix after one hot encoding (109248, 16512)  
16512

### 1.7.3 1.5.3 Vectorizing Numerical features

```

In [0]: # check this one: https://www.youtube.com/watch?v=0HQq0cln3Z4&t=530s
# standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.p
from sklearn.preprocessing import Normalizer

# price_standardized = standardScalar.fit(project_data['price'].values)
# this will rise the error
# ValueError: Expected 2D array, got 1D array instead: array=[725.05 213.03 329. ...
# Reshape your data either using array.reshape(-1, 1)

price_scalar = Normalizer()
price_scalar.fit(project_data['price'].values.reshape(-1,1)) # finding the mean and st

# Now standardize the data with above maen and variance.
price_standardized = price_scalar.transform(project_data['price'].values.reshape(-1, 1))
tr_price_standardized = price_scalar.transform(X_train['price'].values.reshape(-1, 1))
te_price_standardized = price_scalar.transform(X_test['price'].values.reshape(-1, 1))

In [0]: # check this one: https://www.youtube.com/watch?v=0HQq0cln3Z4&t=530s
# standardization sklearn: https://scikit-learn.org/stable/modules/generated/sklearn.p
from sklearn.preprocessing import Normalizer

# price_standardized = standardScalar.fit(project_data['price'].values)
# this will rise the error
# ValueError: Expected 2D array, got 1D array instead: array=[725.05 213.03 329. ...
# Reshape your data either using array.reshape(-1, 1)

quantity_scalar = Normalizer()
quantity_scalar.fit(project_data['quantity'].values.reshape(-1,1)) # finding the mean

# Now standardize the data with above maen and variance.
quantity_standardized = quantity_scalar.transform(project_data['quantity'].values.resha
tr_quantity_standardized = quantity_scalar.transform(X_train['quantity'].values.reshap
te_quantity_standardized = quantity_scalar.transform(X_test['quantity'].values.reshape

```

### 1.3 Make Data Model Ready: encoding eassay, and project\_title

## 2 Encoding - Categorical

```
In [59]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)

         vectorizer = CountVectorizer()
         vectorizer.fit(X_train['clean_categories'].values) # fit has to happen only on train

         # we use the fitted CountVectorizer to convert the text to vector
         X_train_cl_categories_ohe = vectorizer.transform(X_train['clean_categories'].values)
         X_test_cl_categories_ohe = vectorizer.transform(X_test['clean_categories'].values)

         print("After vectorizations")
         print(X_train_cl_categories_ohe.shape, y_train.shape)
         print(X_test_cl_categories_ohe.shape, y_test.shape)
         print("="*100)

         # print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
         # vectorizer = CountVectorizer()
         # x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
         # x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)
         # x_test_bow = vectorizer.fit_transform(X_test['essay'].values)

         # print(x_train_bow.shape, y_train.shape)
         # print(x_cv_bow.shape, y_cv.shape)
         # print(x_test_bow.shape, y_test.shape)

         print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

(87398, 19) (87398,)
(21850, 19) (21850,)
=====
After vectorizations
(87398, 9) (87398,)
(21850, 9) (21850,)
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME

In [60]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)
```

```

vectorizer = CountVectorizer()
vectorizer.fit(X_train['clean_subcategories'].values) # fit has to happen only on tra

# we use the fitted CountVectorizer to convert the text to vector
X_train_cl_subcategories_ohe = vectorizer.transform(X_train['clean_subcategories'].va
X_test_cl_subcategories_ohe = vectorizer.transform(X_test['clean_subcategories'].valu

print("After vectorizations")
print(X_train_cl_subcategories_ohe.shape, y_train.shape)
print(X_test_cl_subcategories_ohe.shape, y_test.shape)
print("="*100)

print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

(87398, 19) (87398,)
(21850, 19) (21850,)
=====
After vectorizations
(87398, 30) (87398,)
(21850, 30) (21850,)
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME

```

```

In [61]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)

vectorizer = CountVectorizer()
vectorizer.fit(X_train['school_state'].values) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_school_state_ohe = vectorizer.transform(X_train['school_state'].values)
X_test_school_state_ohe = vectorizer.transform(X_test['school_state'].values)

print("After vectorizations")
print(X_train_school_state_ohe.shape, y_train.shape)
print(X_test_school_state_ohe.shape, y_test.shape)
print("="*100)

# print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
# vectorizer = CountVectorizer()
# x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
# x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)

```

```

# x_test_bow = vectorizer.fit_transform(X_test['essay'].values)

# print(x_train_bow.shape, y_train.shape)
# print(x_cv_bow.shape, y_cv.shape)
# print(x_test_bow.shape, y_test.shape)

print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

(87398, 19) (87398,)
(21850, 19) (21850,)
=====
After vectorizations
(87398, 51) (87398,)
(21850, 51) (21850,)
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME

In [77]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)

vectorizer = CountVectorizer(vocabulary=list(sorted_fix_dict.keys()), lowercase=False)
vectorizer.fit(X_train['teacher_prefix'].values.astype('U')) # fit has to happen only
#https://stackoverflow.com/questions/39303912/tfidfvectorizer-in-scikit-learn-valueerror
# we use the fitted CountVectorizer to convert the text to vector
X_train_teacher_prefix_ohe = vectorizer.transform(X_train['teacher_prefix'].values.astype('U'))
X_test_teacher_prefix_ohe = vectorizer.transform(X_test['teacher_prefix'].values.astype('U'))

print("After vectorizations")
print(X_train_teacher_prefix_ohe.shape, y_train.shape)
print(X_test_teacher_prefix_ohe.shape, y_test.shape)
print("="*100)

# print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
# vectorizer = CountVectorizer()
# x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
# x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)
# x_test_bow = vectorizer.fit_transform(X_test['essay'].values)

# print(x_train_bow.shape, y_train.shape)
# print(x_cv_bow.shape, y_cv.shape)
# print(x_test_bow.shape, y_test.shape)

print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

```

```
(87398, 19) (87398,)
(21850, 19) (21850,)
```

```
=====
After vectorizations
```

```
(87398, 5) (87398,)
(21850, 5) (21850,)
```

```
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME
```

```
In [69]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)
```

```
print("="*100)
```

```
vectorizer1 = CountVectorizer(vocabulary=list(sorted_sub_cat_dict.keys()), lowercase=1)
vectorizer1.fit(X_train['grade_cat_list'].values) # fit has to happen only on train data
```

```
# we use the fitted CountVectorizer to convert the text to vector
```

```
X_train_grade_ohe = vectorizer1.transform(X_train['grade_cat_list'].values)
```

```
X_test_grade_ohe = vectorizer1.transform(X_test['grade_cat_list'].values)
```

```
print("After vectorizations")
```

```
print(X_train_grade_ohe.shape, y_train.shape)
```

```
print(X_test_grade_ohe.shape, y_test.shape)
```

```
print("="*100)
```

```
# print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
```

```
# vectorizer = CountVectorizer()
```

```
# x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
```

```
# x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)
```

```
# x_test_bow = vectorizer.fit_transform(X_test['essay'].values)
```

```
# print(x_train_bow.shape, y_train.shape)
```

```
# print(x_cv_bow.shape, y_cv.shape)
```

```
# print(x_test_bow.shape, y_test.shape)
```

```
print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")
```

```
(87398, 19) (87398,)
(21850, 19) (21850,)
```

```
=====
After vectorizations
```

```
(87398, 4) (87398,)
(21850, 4) (21850,)
```



NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME

#

2.3 Make Data Model Ready: encoding eassay, and project\_title

### 3 Text - BOW

```
In [70]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)

vectorizer = CountVectorizer(min_df=10, max_features=5000)
vectorizer.fit(train_preprocessed_essays) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_essay_bow = vectorizer.transform(train_preprocessed_essays)
X_test_essay_bow = vectorizer.transform(test_preprocessed_essays)

print("After vectorizations")
print(X_train_essay_bow.shape, y_train.shape)
print(X_test_essay_bow.shape, y_test.shape)
print("="*100)

# print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
# vectorizer = CountVectorizer()
# x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
# x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)
# x_test_bow = vectorizer.fit_transform(X_test['essay'].values)

# print(x_train_bow.shape, y_train.shape)
# print(x_cv_bow.shape, y_cv.shape)
# print(x_test_bow.shape, y_test.shape)

print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

(87398, 19) (87398,)
(21850, 19) (21850,)
=====
After vectorizations
(87398, 5000) (87398,)
(21850, 5000) (21850,)
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME
```

## 4 Text - TfIdf

```
In [71]: print(X_train.shape, y_train.shape)
         print(X_test.shape, y_test.shape)

         print("="*100)

vectorizer = TfidfVectorizer(min_df=10, max_features=5000)
vectorizer.fit(train_preprocessed_essays) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_essay_tf = vectorizer.transform(train_preprocessed_essays)
X_test_essay_tf = vectorizer.transform(test_preprocessed_essays)

print("After vectorizations")
print(X_train_essay_tf.shape, y_train.shape)
print(X_test_essay_tf.shape, y_test.shape)
print("="*100)

# print("YOU SHOULD NOT DO SOMETHING LIKE THIS")
# vectorizer = CountVectorizer()
# x_train_bow = vectorizer.fit_transform(X_train['essay'].values)
# x_cv_bow = vectorizer.fit_transform(X_cv['essay'].values)
# x_test_bow = vectorizer.fit_transform(X_test['essay'].values)

# print(x_train_bow.shape, y_train.shape)
# print(x_cv_bow.shape, y_cv.shape)
# print(x_test_bow.shape, y_test.shape)

print("NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME")

(87398, 19) (87398,)
(21850, 19) (21850,)
=====
After vectorizations
(87398, 5000) (87398,)
(21850, 5000) (21850,)
=====
NOTE: THE NUMBER OF COLUMNS IN EACH OF THE VECTOR WONT BE SAME
```

1.5 Appling NB on different kind of featurization as mentioned in the instructions  
Apply NB on different kind of featurization as mentioned in the instructions For Every model  
that you work on make sure you do the step 2 and step 3 of instrucations

## 5 Concatinating all the features

```
In [81]: from scipy.sparse import hstack
         # with the same hstack function we are concatinating a sparse matrix and a dense matrix
         X_Bow_train = hstack((X_train_essay_bow, tr_price_standardized, tr_quantity_standardized,
                               X_train_cl_subcategories_ohe, X_train_school_state_ohe, X_train_teachers_ohe))

         X_Bow_train=X_Bow_train.tocsr()
         print(X_Bow_train.shape, y_train.shape)

(87398, 5101) (87398,)
```

```
In [82]: X_Bow_test = hstack((X_test_essay_bow, te_price_standardized, te_quantity_standardized,
                              X_test_cl_subcategories_ohe, X_test_school_state_ohe, X_test_teachers_ohe))

         X_Bow_test = X_Bow_test.tocsr()
         print(X_Bow_test.shape, y_test.shape)

(21850, 5101) (21850,)
```

```
In [0]: #TfIdf
```

```
In [83]: from scipy.sparse import hstack
         # with the same hstack function we are concatinating a sparse matrix and a dense matrix
         X_tf_train = hstack((X_train_essay_tf, tr_price_standardized, tr_quantity_standardized,
                               X_train_cl_subcategories_ohe, X_train_school_state_ohe, X_train_teachers_ohe))

         X_tf_train = X_tf_train.tocsr()
         print(X_tf_train.shape, y_train.shape)

(87398, 5101) (87398,)
```

```
In [84]: X_tf_test = hstack((X_test_essay_tf, te_quantity_standardized, te_price_standardized,
                              X_test_cl_subcategories_ohe, X_test_school_state_ohe, X_test_teachers_ohe))

         X_tf_test = X_tf_test.tocsr()
         print(X_tf_test.shape, y_test.shape)

(21850, 5101) (21850,)
```

## 6 Set: 1

```
In [85]: # https://scikit-learn.org/stable/modules/generated/sklearn.model\_selection.GridSearchCV
         from sklearn.model_selection import GridSearchCV
```

```

from scipy.stats import randint as sp_randint
from sklearn.model_selection import RandomizedSearchCV
from sklearn.naive_bayes import MultinomialNB

alpha=[0.0001,0.0001,0.001,0.1,1,10,100,1000,10000]
MNB = MultinomialNB()
parameters = dict(alpha=[0.0001,0.0001,0.001,0.1,1,10,100,1000,10000])
clf = GridSearchCV(MNB, parameters, scoring='roc_auc', return_train_score=True)
clf.fit(X_Bow_train, y_train)

results = pd.DataFrame.from_dict(clf.cv_results_)
results = results.sort_values(['param_alpha'])

train_auc= results['mean_train_score']
train_auc_std= results['std_train_score']
cv_auc = results['mean_test_score']
cv_auc_std= results['std_test_score']
#K = results['param_alpha']

plt.plot(alpha, train_auc, label='Train AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, train_auc - train_auc_std, train_auc + train_auc_std, alpha=0.2, color='r')


plt.plot(alpha, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.2, color='r')

plt.scatter(alpha, train_auc, label='Train AUC points')
plt.scatter(alpha, cv_auc, label='CV AUC points')

plt.xscale('log')# we take the log in the x axis
plt.legend()
plt.xlabel("Alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("Hyper parameter Vs AUC plot")
plt.grid()
plt.show()

results.head()

```



```
Out [85]:
```

	mean_fit_time	std_fit_time	...	mean_train_score	std_train_score
0	0.079387	0.001326	...	0.733684	0.000650
1	0.079662	0.000581	...	0.733684	0.000650
2	0.079192	0.001937	...	0.733678	0.000653
3	0.078312	0.000674	...	0.733472	0.000670
4	0.078341	0.000758	...	0.731736	0.000704

[5 rows x 21 columns]

```
In [86]: #https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearch
print(clf.best_estimator_)
print(clf.best_index_)
a=clf.best_params_["alpha"]
print("Best Parameter Found:- ",a)
print(clf.best_score_)
```

```
MultinomialNB(alpha=0.001, class_prior=None, fit_prior=True)
2
Best Parameter Found:- 0.001
0.7013408748453728
```

```
In [87]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sk
from sklearn.metrics import roc_curve, auc
```

```
MNB = MultinomialNB(alpha=a) # n_jobs=-1 means parallel operations
MNB.fit(X_Bow_train, y_train)
```

```
Out[87]: MultinomialNB(alpha=0.001, class_prior=None, fit_prior=True)
```

```
In [88]: from sklearn.metrics import roc_curve
from sklearn.metrics import auc
import matplotlib.pyplot as plt
```


```
score_roc_train = MNB.predict_proba(X_Bow_train)
fpr_train, tpr_train, threshold_train = roc_curve(y_train, score_roc_train[:,1])
roc_auc_train = auc(fpr_train, tpr_train)
```

```
score_roc_test = MNB.predict_proba(X_Bow_test)
fpr_test, tpr_test, threshold_test = roc_curve(y_test, score_roc_test[:,1])
roc_auc_test = auc(fpr_test, tpr_test)
```

```
plt.plot(fpr_train, tpr_train, label = "Train_AUC"+str(auc(fpr_train, tpr_train)))
plt.plot(fpr_test, tpr_test, label = "Test_AUC"+str(auc(fpr_test, tpr_test)))
plt.legend(loc = 'lower right')
```


```
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
```

```
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.title('ROC Curve of NB ')
plt.show()
```



```
In [89]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
import seaborn as sns
import matplotlib.pyplot as plt
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_train, MNB.predict(X_Bow_train)), annot=True, ax = ax,


# labels, title and ticks
ax.set_xlabel('Predicted labels');
ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
```



```
In [90]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
#Test Confusion Matrix
import seaborn as sns
import matplotlib.pyplot as plt
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test, MNB.predict(X_Bow_test)), annot=True, ax = ax,fmt

# labels, title and ticks
ax.set_xlabel('Predicted labels');
ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
```





## 7 Set: 2

```
In [91]: # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
from sklearn.model_selection import GridSearchCV
from scipy.stats import randint as sp_randint
from sklearn.model_selection import RandomizedSearchCV

alpha=[0.0001,0.0001,0.001,0.1,1,10,100,1000,10000]
MNB1 = MultinomialNB()
parameters = dict(alpha=[0.0001,0.0001,0.001,0.1,1,10,100,1000,10000])
clf1 = GridSearchCV(MNB1, parameters, scoring='roc_auc',return_train_score=True)
clf1.fit(X_tf_train, y_train)

results = pd.DataFrame.from_dict(clf1.cv_results_)
results = results.sort_values(['param_alpha'])

train_auc= results['mean_train_score']
train_auc_std= results['std_train_score']
cv_auc = results['mean_test_score']
cv_auc_std= results['std_test_score']
#K = results['param_alpha']

plt.plot(alpha, train_auc, label='Train AUC')
```

```

# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, train_auc - train_auc_std, train_auc + train_auc_std, alpha=0.2, color='blue')


plt.plot(alpha, cv_auc, label='CV AUC')
# this code is copied from here: https://stackoverflow.com/a/48803361/4084039
# plt.gca().fill_between(K, cv_auc - cv_auc_std, cv_auc + cv_auc_std, alpha=0.2, color='orange')

plt.scatter(alpha, train_auc, label='Train AUC points')
plt.scatter(alpha, cv_auc, label='CV AUC points')

plt.xscale('log') # we take the log in the x axis
plt.legend()
plt.xlabel("Alpha: hyperparameter")
plt.ylabel("AUC")
plt.title("Hyper parameter Vs AUC plot")
plt.grid()
plt.show()

results.head()

```



```

Out[91]:
  mean_fit_time  std_fit_time  ...  mean_train_score  std_train_score
0      0.080035      0.005750  ...           0.710610          0.000673
1      0.075708      0.000932  ...           0.710610          0.000673

```

2	0.075102	0.000728	...	0.710589	0.000675
3	0.075453	0.000860	...	0.708872	0.000703
4	0.074701	0.000257	...	0.693549	0.000760

[5 rows x 21 columns]

```
In [92]: #https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV
print(clf.best_estimator_)
print(clf.best_index_)
a1=clf.best_params_["alpha"]
print("Best Parameter Found:- ",clf.best_params_)
print(clf.best_score_)
```

```
MultinomialNB(alpha=0.001, class_prior=None, fit_prior=True)
```

```
2
```

```
Best Parameter Found:- {'alpha': 0.001}
```

```
0.7013408748453728
```

```
In [93]: # https://scikit-learn.org/stable/modules/generated/sklearn.metrics.roc_curve.html#sk
from sklearn.metrics import roc_curve, auc
MNB1 = MultinomialNB(alpha=a1) # n_jobs=-1 means parallel operations
MNB1.fit(X_tf_train, y_train)
```

```
Out[93]: MultinomialNB(alpha=0.001, class_prior=None, fit_prior=True)
```

```
In [94]: from sklearn.metrics import roc_curve
from sklearn.metrics import auc
import matplotlib.pyplot as plt
```

```
score_roc_train = MNB1.predict_proba(X_tf_train)
fpr_train, tpr_train, threshold_train = roc_curve(y_train, score_roc_train[:,1])
roc_auc_train = auc(fpr_train, tpr_train)
```

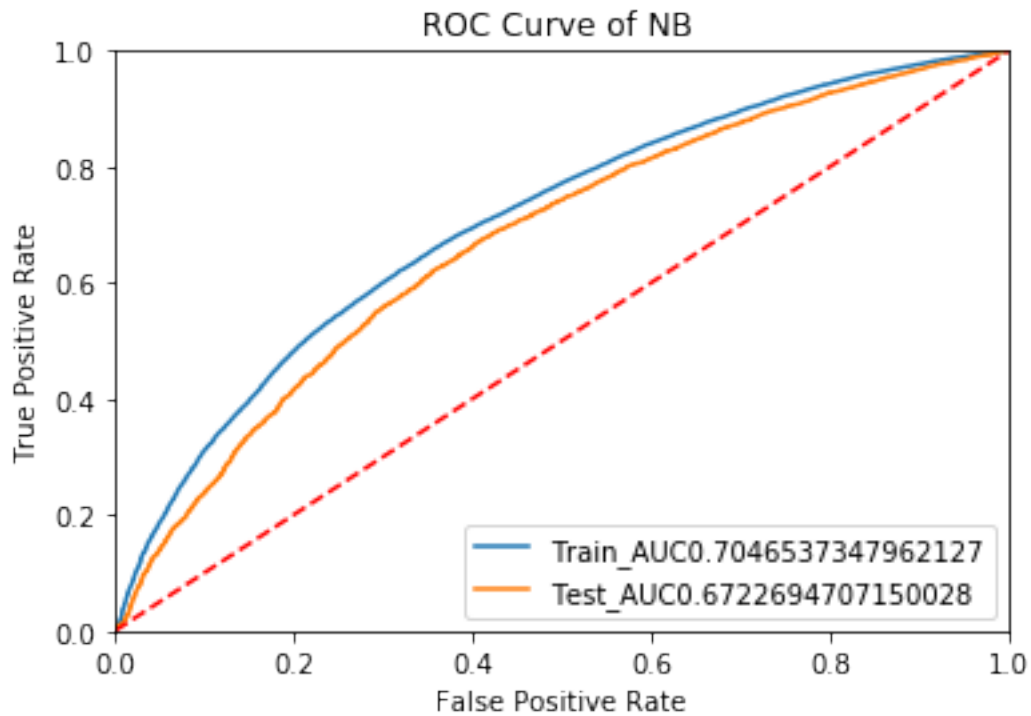
```
score_roc_test = MNB1.predict_proba(X_tf_test)
fpr_test, tpr_test, threshold_test = roc_curve(y_test, score_roc_test[:,1])
roc_auc_test = auc(fpr_test, tpr_test)
```

```
plt.plot(fpr_train, tpr_train, label = "Train_AUC"+str(auc(fpr_train, tpr_train)))
plt.plot(fpr_test, tpr_test, label = "Test_AUC"+str(auc(fpr_test, tpr_test)))
plt.legend(loc = 'lower right')
```

```
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
```


```
plt.ylabel('True Positive Rate')
```

```
plt.xlabel('False Positive Rate')
plt.title('ROC Curve of NB ')
plt.show()
```




```
In [95]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
import seaborn as sns
import matplotlib.pyplot as plt
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_train, MNB.predict(X_tf_train)), annot=True, ax = ax,

# labels, title and ticks
ax.set_xlabel('Predicted labels');
ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
```



```
In [96]: #https://stackoverflow.com/questions/35572000/how-can-i-plot-a-confusion-matrix
import seaborn as sns
import matplotlib.pyplot as plt
ax= plt.subplot()
sns.heatmap(confusion_matrix(y_test, MNB.predict(X_tf_test)), annot=True, ax = ax,fmt=

# labels, title and ticks
ax.set_xlabel('Predicted labels');
ax.set_ylabel('Actual labels');
ax.set_title('Confusion Matrix');
```



## 8 Top 10 features (negatives and positives)

## 9 Applying top 10 features on BoW

In [0]: <https://stackoverflow.com/questions/54988116/why-am-i-getting-almost-same-top-10-feat>

In [0]: essays\_features\_tf=vectorizer3.get\_feature\_names()

In [100]: sorted\_idx = np.argsort(MNB.feature\_log\_prob\_[1] )[-10:]

```
for i in sorted_idx:
    print(vectorizer3.get_feature_names()[i])
```

```
cliche
embracing
embraces
busy
chew
complexity
always
chicagoland
defensive
discouraging
```

```
In [101]: sorted_idx = np.argsort(-1 * MNB.feature_log_prob_[0] )[0:11]

        for i in sorted_idx:
            print(vectorizer3.get_feature_names()[i])
```

```
discouraging
defensive
chicagoland
always
complexity
chew
busy
embracing
embraces
commonalities
cliche
```

## 10 Conclusion

```
In [102]: # Please compare all your models using Prettytable library
          #how to use pretty table http://zetcode.com/python/prettytable/
          from prettytable import PrettyTable
          tb = PrettyTable()
          tb.field_names= ("Vectorizer", "HyperParameter", "AUC")
          tb.add_row(["BOW", a, 0.699])
          tb.add_row(["Tf-Idf", a1 , 0.672])
          print(tb.get_string(titles = "Naive Bayes - Observations"))
          #print(tb)
```

```
+-----+-----+-----+
| Vectorizer | HyperParameter | AUC |
+-----+-----+-----+
| BOW       | 0.001         | 0.699 |
| Tf-Idf    | 0.001         | 0.672 |
+-----+-----+-----+
```